

# Unleashing Transformers: Parallel Token Prediction with Discrete Absorbing Diffusion for Fast High-Resolution Image Generation from Vector-Quantized Codes

Sam Bond-Taylor\*, Peter Hessey\*, Hiroshi Sasaki, Toby P. Breckon, and Chris G. Willcocks



Project Page

## Introduction

A parallel prediction approach that allows **fast generation** of high-res Vector-Quantized images:

- SOTA quantitative sample quality scores.
- Sample resolutions higher than training data.
- Global context allows samples to be edited.



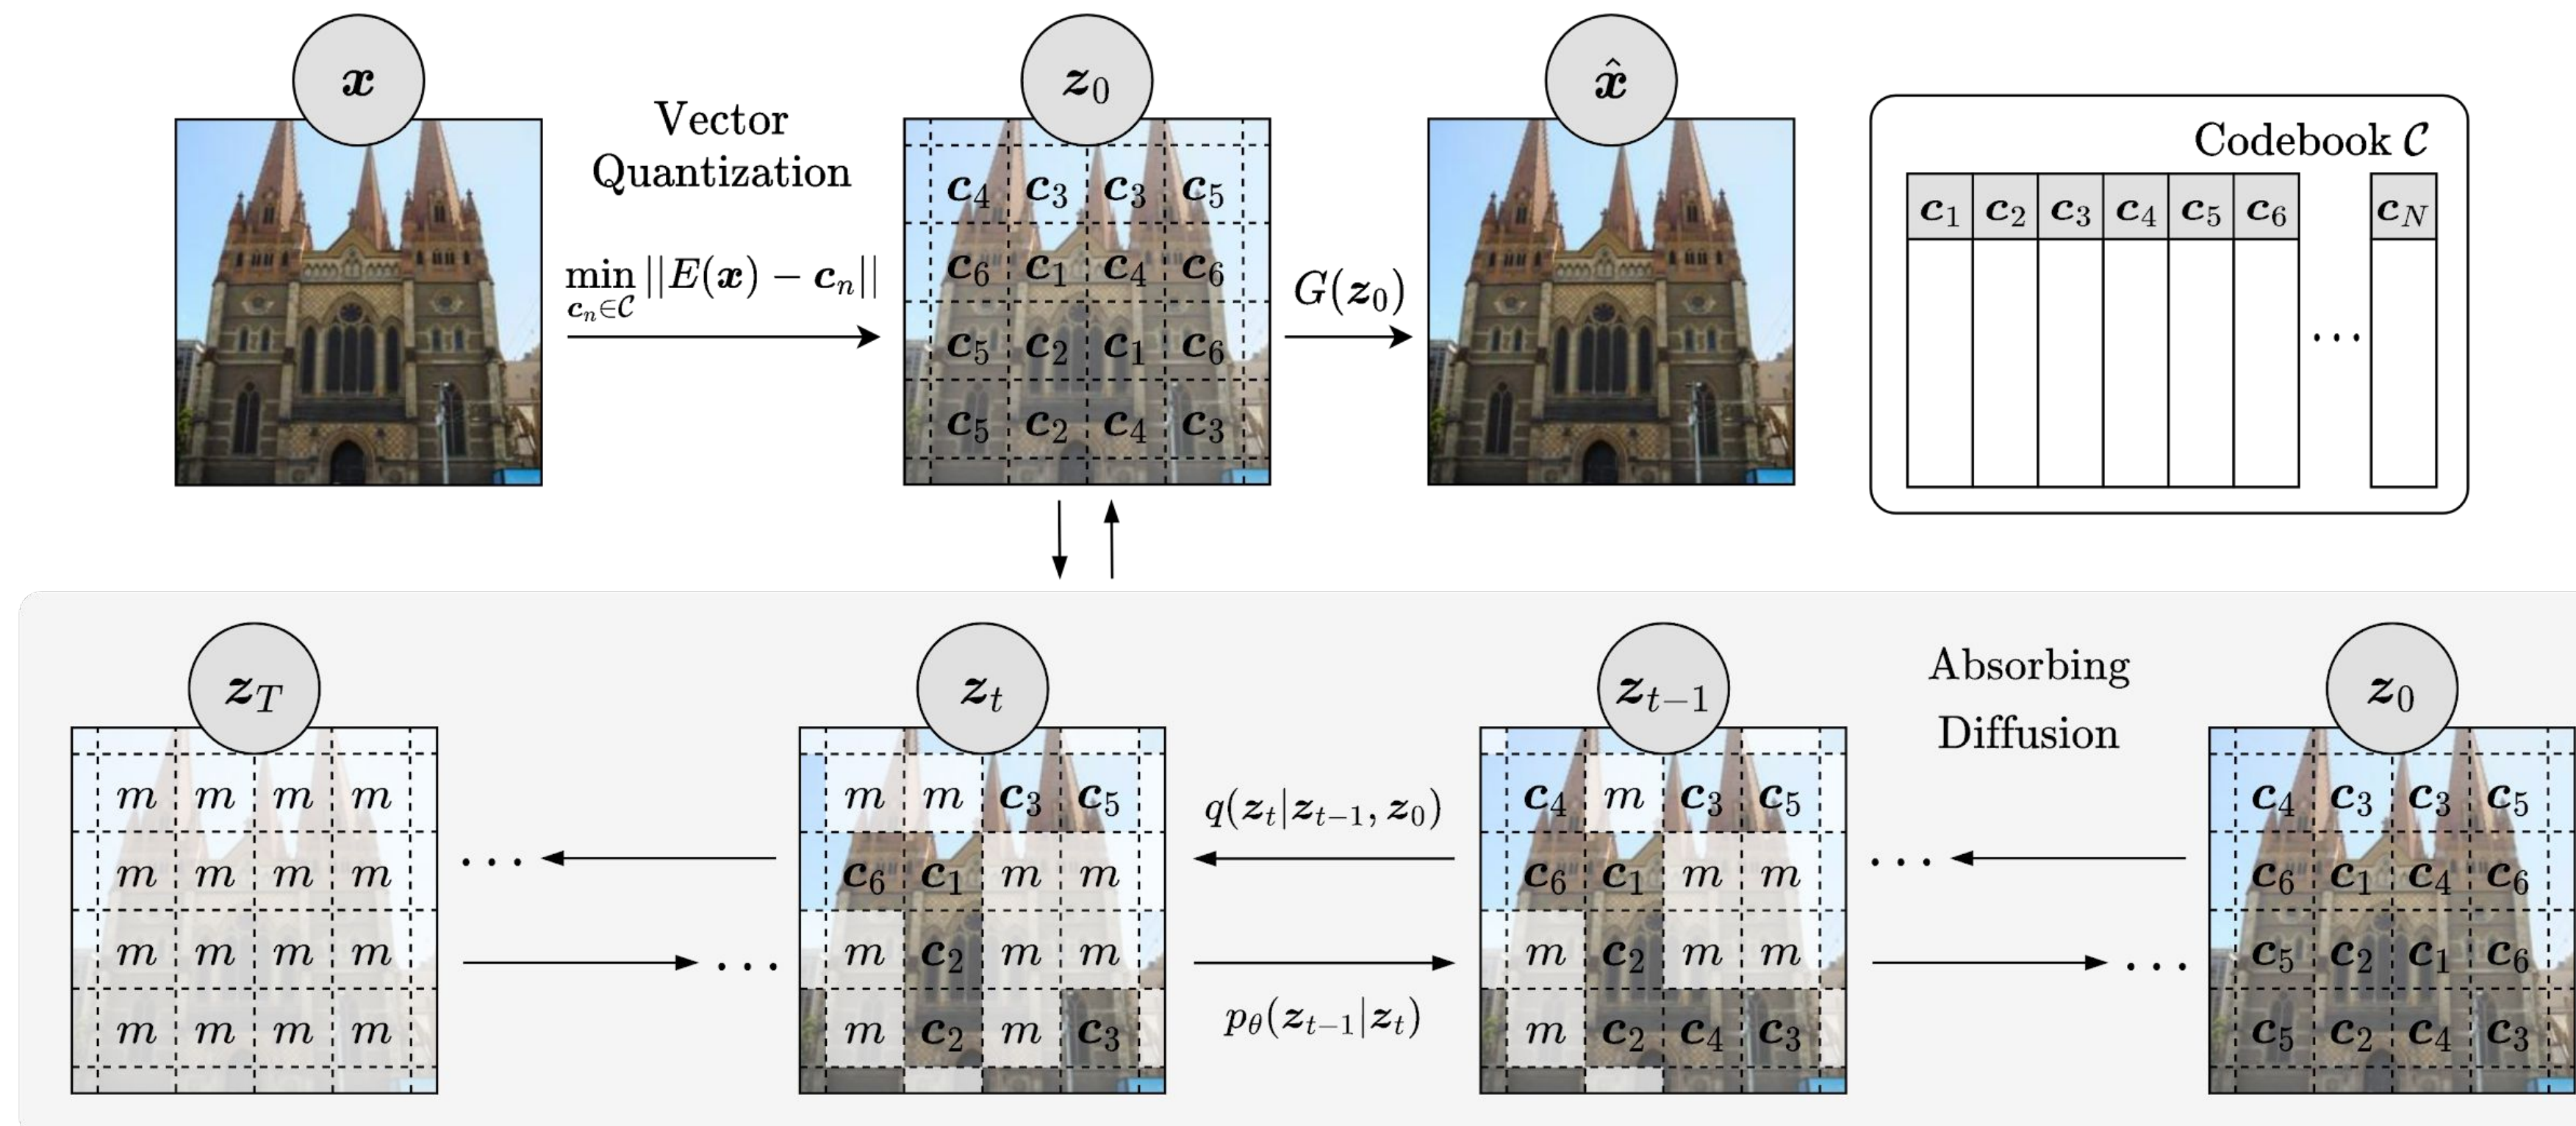
## Method

Discrete latents  $\mathbf{z}$  obtained using a VQ-VAE are modelled with discrete **absorbing diffusion**:

$$p_{\theta}(\mathbf{z}_{0:T}) = p_{\theta}(\mathbf{z}_T) \prod_{t=1}^T p_{\theta}(\mathbf{z}_{t-1} | \mathbf{z}_t)$$

Elements of  $\mathbf{z}$  are randomly masked and an **unconstrained Transformer** learns to denoise the data by optimising a carefully reweighted ELBO designed to improve convergence rates:

$$\sum_{t=1}^T \frac{T-t+1}{T} \mathbb{E}_{q(\mathbf{z}_t | \mathbf{z}_0)} \left[ \sum_{[z_t]_i=m} \log p_{\theta}([z_0]_i | \mathbf{z}_t) \right]$$



A discrete absorbing diffusion process destroys latents by slowly masking out tokens over many steps, similar to masked language models like BERT.


## Sampling at Higher Resolutions

Globally consistent images at **higher resolutions** than the training data can be generated by aggregating multiple context windows:



## Image Editing

Our bidirectional approach with **global context** allows internal image regions to be edited by masking those areas (highlighted in grey).



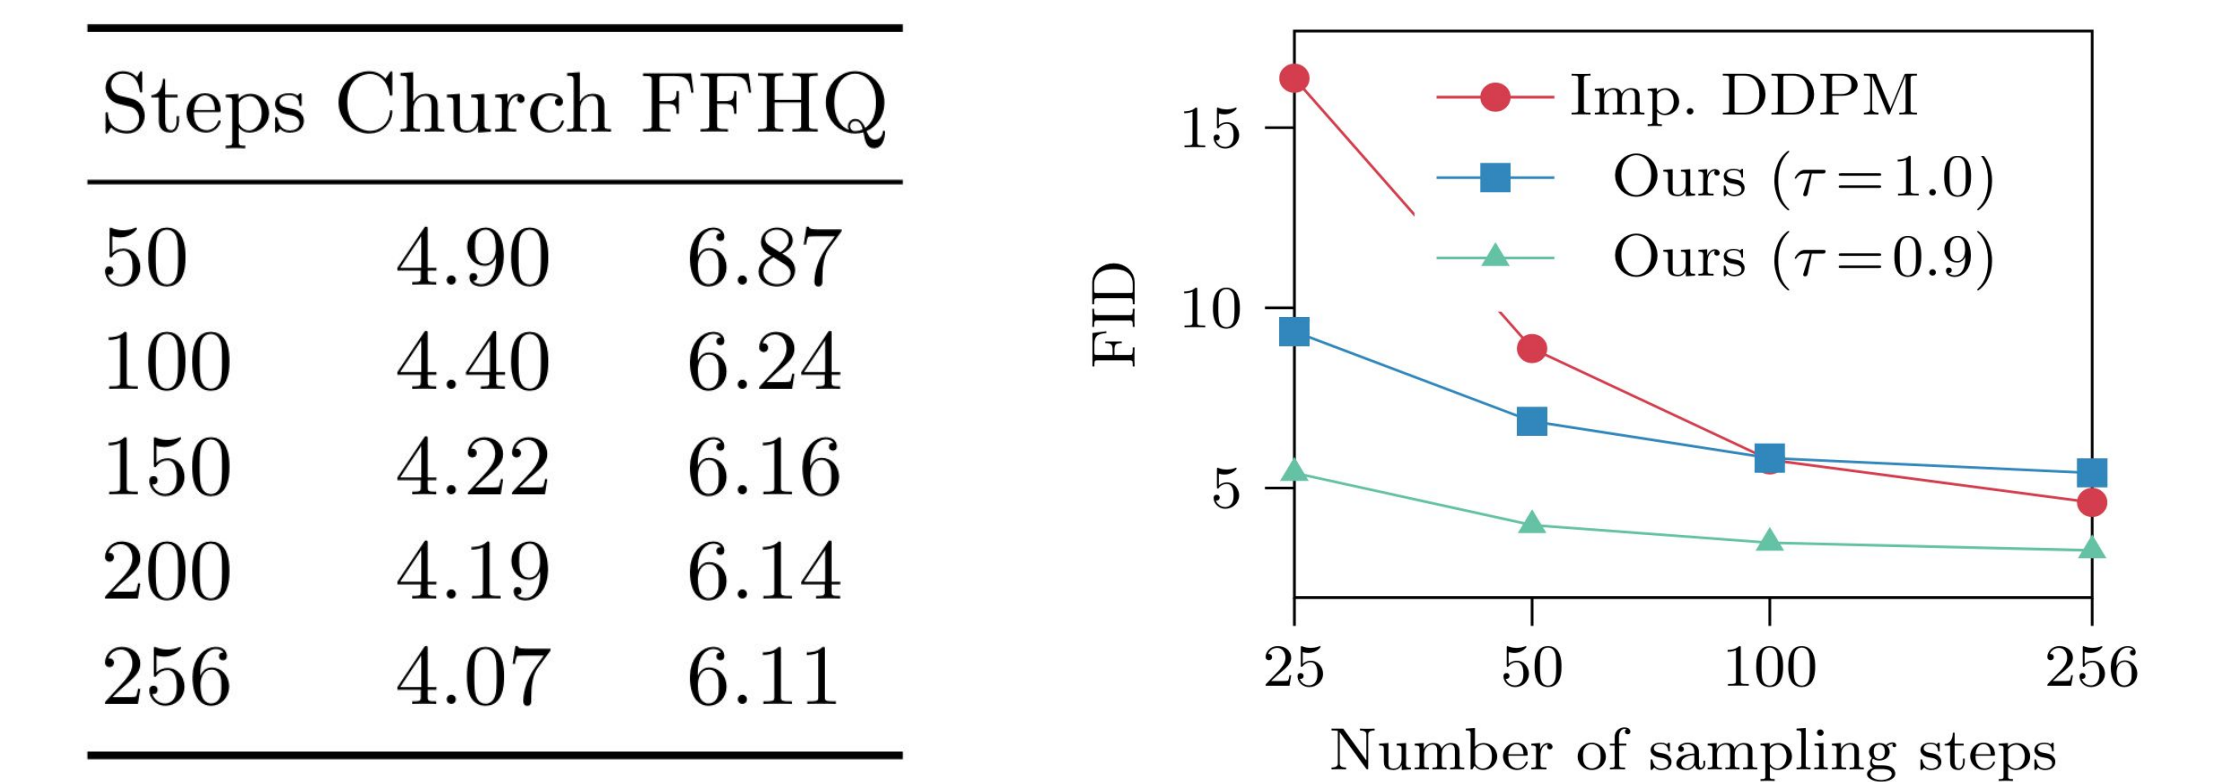
## Quantitative Results

Our models achieve competitive FID scores on LSUN and FFHQ. Evaluating our models using Precision, Recall, Density and Coverage (PRDC) metrics further demonstrates **SOTA** results:

Model	Churches			Bedroom			FFHQ		
	FID ↓	D ↑	C ↑	FID ↓	D ↑	C ↑	FID ↓	D ↑	C ↑
TT	7.81	<b>1.08</b>	0.60	6.35	1.15	0.75	9.6	0.89	0.50
ImageBART	7.32	-	-	5.51	-	-	9.57	-	-
StyleGAN2	3.85	0.83	0.68	2.35	-	-	3.80	1.12	0.80
ProjGAN	<b>1.59</b>	0.65	0.64	<b>1.52</b>	0.90	0.79	<b>3.39</b>	0.98	0.77
<b>Ours</b>	4.07	1.07	<b>0.74</b>	3.27	<b>1.51</b>	<b>0.83</b>	6.11	<b>1.20</b>	<b>0.80</b>

## Sampling Speed

**Faster sampling** can be achieved by predicting tokens in parallel with only small FID change:



## Summary

Using a discrete absorbing diffusion model parameterised by an unconstrained Transformer to model VQ-VAE representations we achieve faster sampling with higher visual quality.

**Github** repository with trained models at <https://samb-t.github.io/unleashing-transformers>  
\*Authors contributed equally